

Comparing BM25, Dense Retrieval, and Hybrid Re-ranking

Prashanth Kumar (pk3047)
Pranav Joshi (pj2490)

Abstract

This report presents a comprehensive comparison of three information retrieval systems: (1) BM25 sparse retrieval, (2) HNSW-based dense vector retrieval, and (3) a hybrid re-ranking system combining both approaches. We implemented all three systems on a subset of 1 million passages from the MS-MARCO dataset and evaluated them on TREC Deep Learning 2019/2020 benchmarks. Our results demonstrate that while BM25 achieves solid baseline performance (MRR@10: 56-59%), dense retrieval excels at recall (93.21%), and the hybrid system delivers the best overall performance (MRR@10: 82-90%, NDCG@10: 48-59%, NDCG@100: 43-48%). The hybrid approach improves MRR@10 by 47-112% over BM25 across all test sets. We analyze the trade-offs in retrieval quality, computational efficiency, memory usage, and provide insights into when each approach is most suitable.

Contents

1	Introduction	5
1.1	Background and Motivation	5
1.2	Dataset and Experimental Setup	5
2	System Architecture	5
2.1	Project Implementation Files	6
2.2	Execution Workflow	7
2.2.1	BM25 System	7
2.2.2	HNSW Dense Retrieval	7
2.2.3	Hybrid Re-ranking	8
2.3	System Components	8
2.3.1	Input Data Layer	8
2.3.2	Preprocessing Layer	8
2.3.3	Indexing Layer	8
3	Implementation Details	9
3.1	System 1: BM25 Sparse Retrieval	9
3.1.1	Algorithm	9
3.1.2	Implementation	9
3.1.3	Performance Characteristics	9
3.2	System 2: HNSW Dense Vector Retrieval	10
3.2.1	Algorithm	10
3.2.2	Implementation	10
3.2.3	Performance Characteristics	10
3.3	System 3: Hybrid Re-ranking	11
3.3.1	Algorithm	11
3.3.2	Implementation	11
3.3.3	Performance Characteristics	11
4	Evaluation Results	12
4.1	TREC Deep Learning 2019 (qrels.eval.one.tsv)	12
4.2	TREC Deep Learning 2020 (qrels.eval.two.tsv)	12
4.3	Dev Set (qrels.dev.tsv)	12
4.4	Key Observations	12
5	Analysis and Discussion	13
5.1	Trade-offs in Retrieval Quality	13
5.1.1	Precision vs. Recall	13
5.1.2	When Each System Excels	14
5.2	Efficiency Analysis	14
5.2.1	Time Complexity	14
5.2.2	Memory Usage	15
5.2.3	Query Latency	15

6 Conclusion	15
6.1 Summary of Findings	15
6.2 Summary	16
6.3 Final Remarks	16

1 Introduction

1.1 Background and Motivation

Information retrieval systems face a fundamental challenge: matching user queries with relevant documents when the query terms may not exactly match the document terms. Traditional sparse retrieval methods like BM25 excel at exact term matching but struggle with semantic similarity. Dense retrieval methods using pre-trained language models address this limitation by encoding queries and documents as high-dimensional vectors, enabling semantic matching through vector similarity.

This project implements and compares three retrieval paradigms:

1. **BM25 Sparse Retrieval:** A probabilistic ranking function using inverted indexes with TF-IDF weighting
2. **HNSW Dense Retrieval:** Graph-based approximate nearest neighbor search using pre-computed 384-dimensional embeddings
3. **Hybrid Re-ranking:** BM25 candidate generation followed by dense vector re-ranking

1.2 Dataset and Experimental Setup

We use a curated subset of the MS-MARCO passage ranking dataset:

- **Passages:** 1 million passages (from 8.8M total)
- **Embeddings:** 384-dimensional vectors generated using MiniLM
- **Similarity:** Dot product (normalized to cosine similarity)
- **Evaluation Sets:**
 - TREC DL 2019 (qrels.eval.one.tsv): 9,260 graded judgments (0-3)
 - TREC DL 2020 (qrels.eval.two.tsv): 11,385 graded judgments (0-3)
 - Dev Set (qrels.dev.tsv): 1,073 binary judgments (0-1)
- **Metrics:** MRR@10, Recall@100, NDCG@10, NDCG@100, MAP

2 System Architecture

Figure 1 illustrates the complete system architecture showing the data flow through all three retrieval systems and the evaluation pipeline.

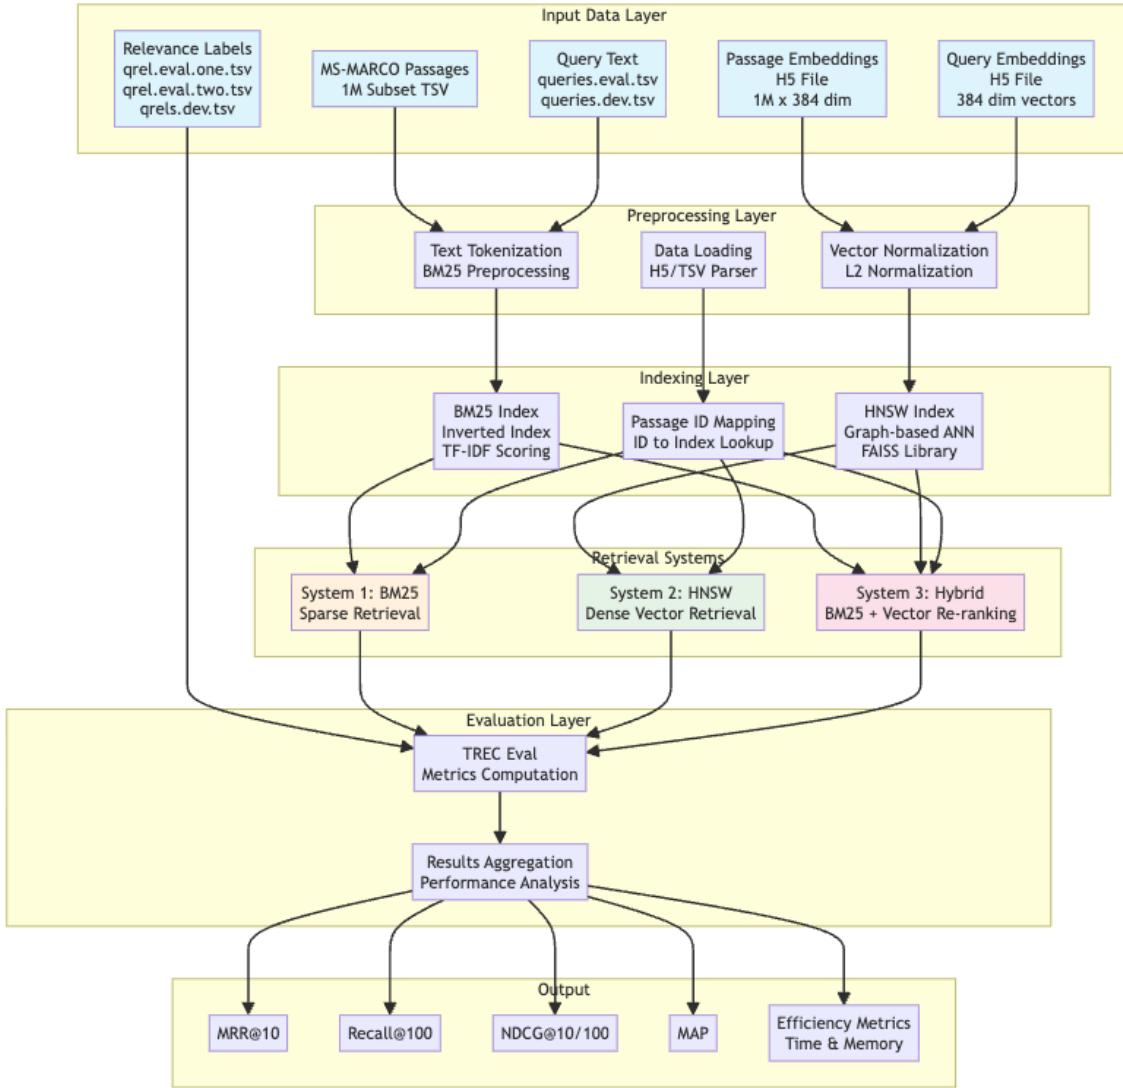


Figure 1: Complete system architecture showing BM25, HNSW, and Hybrid re-ranking systems

2.1 Project Implementation Files

The project consists of the following source code files organized by system:

Table 1: Source code organization

File	Description
<i>BM25 System (C++)</i>	
BM25/indexer.cpp	Builds inverted index from passages using SPIMI algorithm
BM25/merger.cpp	Merges partial indexes and applies varbyte compression
BM25/query.cpp	Processes queries and returns top-1000 results using BM25
<i>Dense Retrieval (Python)</i>	
ann.py	Performs approximate nearest neighbor search using HNSW
<i>Hybrid Re-ranking (Python)</i>	
rerank.py	Re-ranks BM25 candidates using dense vector similarity
<i>Utilities</i>	
fix_qrels.py	Converts qrels files to TREC evaluation format
trec_eval	Standard TREC evaluation tool for computing metrics

2.2 Execution Workflow

The complete pipeline for running all three systems:

2.2.1 BM25 System

```
# Compile
g++ -O3 -std=c++17 indexer.cpp -o indexer
g++ -O3 -std=c++17 merger.cpp -o merger
g++ -O3 -std=c++17 query.cpp -o query

# Build index
cd BM25
./indexer
./merger

# Run queries
./query ../queries/queries.eval.tsv
./query ../queries/queries.dev.tsv
```

2.2.2 HNSW Dense Retrieval

```
# Run HNSW search
python3 ann.py

# Output: run_hnsw.txt
```

2.2.3 Hybrid Re-ranking

```
# Run hybrid re-ranking (requires BM25 results)
python3 rerank.py

# Output:
# queries.dev_results_hybrid_hnsw.txt
# queries.eval_results_hybrid_hnsw.txt
```

2.3 System Components

2.3.1 Input Data Layer

- MS-MARCO passages (1M subset TSV)
- Query text files (queries.eval.tsv, queries.dev.tsv)
- Pre-computed passage embeddings (H5 format, 1M × 384 dimensions)
- Pre-computed query embeddings (H5 format, 384 dimensions)
- Relevance labels (qrels files with graded judgments)

2.3.2 Preprocessing Layer

- **BM25 Preprocessing:** Tokenization, lowercase normalization
- **Dense Retrieval:** H5/TSV parsing, L2 normalization for cosine similarity

2.3.3 Indexing Layer

- **BM25 Index:** Inverted index with delta-varbyte compression (98MB)
- **HNSW Index:** Graph-based ANN index using FAISS (in-memory)
- **Passage ID Mapping:** Internal to external ID lookup tables

3 Implementation Details

3.1 System 1: BM25 Sparse Retrieval

3.1.1 Algorithm

BM25 uses the following ranking formula:

$$\text{score}(D, q) = \sum_{t \in q} \text{IDF}(t) \cdot \frac{f(t, D) \cdot (k_1 + 1)}{f(t, D) + k_1 \cdot \left(1 - b + b \cdot \frac{|D|}{\text{avgdl}}\right)} \quad (1)$$

where:

- $f(t, D)$ = term frequency of t in document D
- $|D|$ = document length
- avgdl = average document length
- $k_1 = 1.2, b = 0.75$ (standard parameters)
- $\text{IDF}(t) = \log \left(\frac{N - n(t) + 0.5}{n(t) + 0.5} \right)$

3.1.2 Implementation

- **Language:** C++17
- **Indexing:** SPIMI algorithm with 5 partial runs, k-way merge
- **Compression:** Delta encoding + Varbyte encoding
- **Block Structure:** 128 postings per block for efficient skipping
- **Query Processing:** Multi-threaded (hardware concurrency)
- **Output:** Top 1000 results per query in TREC format

3.1.3 Performance Characteristics

- Index size: 98MB (compressed from 650MB+ raw postings)
- Query latency: 27-58ms for typical queries
- Memory footprint: 195MB (lexicon + metadata loaded in RAM)

3.2 System 2: HNSW Dense Vector Retrieval

3.2.1 Algorithm

Dense retrieval uses cosine similarity between normalized embeddings:

$$\text{score}(q, d) = \cos(q, d) = \frac{q \cdot d}{\|q\| \|d\|} = q \cdot d \quad (\text{after L2 normalization}) \quad (2)$$

HNSW (Hierarchical Navigable Small World) constructs a multi-layer graph where:

- Each layer is a proximity graph with edges to nearest neighbors
- Upper layers have fewer nodes for fast global search
- Lower layers are denser for precise local search
- Search complexity: $O(\log N)$ on average

3.2.2 Implementation

- **Library:** FAISS (Facebook AI Similarity Search)
- **Index Type:** IndexHNSWFlat
- **Parameters:**
 - M = 16 (connections per node)
 - efConstruction = 200 (construction-time search depth)
 - efSearch = 256 (query-time search depth)
- **Preprocessing:** L2 normalization of all vectors
- **Batch Processing:** 100 queries per batch for memory efficiency
- **Output:** Top 100 results per query in TREC format

3.2.3 Performance Characteristics

- Index size: 771MB (passage embeddings in memory)
- Query latency: ~ 2.5 seconds for 202K queries (batch mode)
- Memory footprint: 927MB (771MB embeddings + 156MB query embeddings)

3.3 System 3: Hybrid Re-ranking

3.3.1 Algorithm

The hybrid system combines the strengths of both approaches:

1. **Stage 1 - Candidate Generation:** Use BM25 to retrieve top- K candidates ($K=1000$)
2. **Stage 2 - Re-ranking:** Load embeddings for candidates and compute cosine similarity with query
3. **Stage 3 - Final Ranking:** Sort by similarity and return top 100

This two-stage approach achieves:

- Fast candidate generation (BM25 is extremely efficient)
- Semantic re-ranking (dense vectors capture meaning)
- Scalability (only K embeddings loaded per query instead of all 1M)

3.3.2 Implementation

- **Language:** Python 3.13
- **Libraries:** h5py, numpy, faiss
- **Workflow:**
 1. Load BM25 results (top 1000 per query)
 2. Create passage ID to embedding index mapping
 3. For each query:
 - Extract candidate passage embeddings
 - Compute cosine similarity using FAISS IndexFlatIP
 - Sort by similarity (descending)
 - Output top 100

3.3.3 Performance Characteristics

- Processing time: ~2 minutes for 202K queries
- Memory: 927MB (embeddings) + overhead
- Disk I/O: Minimal (BM25 results pre-computed)

4 Evaluation Results

4.1 TREC Deep Learning 2019 (qrels.eval.one.tsv)

Table 2: Performance on TREC DL 2019 (9,260 graded judgments)

System	MRR@10	NDCG@10	NDCG@100	Recall@100
BM25	0.5925	0.3282	0.3505	0.3386
HNSW Dense	0.2894	0.1464	0.4244	0.5549
Hybrid (top 1000)	0.8956	0.5940	0.4798	0.3897
Improvement	+51.2%	+80.9%	+36.9%	+15.1%

4.2 TREC Deep Learning 2020 (qrels.eval.two.tsv)

Table 3: Performance on TREC DL 2020 (11,385 graded judgments)

System	MRR@10	NDCG@10	NDCG@100	Recall@100
BM25	0.5628	0.2784	0.3222	0.3626
HNSW Dense	0.2669	0.0914	0.3638	0.6129
Hybrid (top 1000)	0.8260	0.4809	0.4300	0.3953
Improvement	+46.8%	+72.7%	+33.5%	+9.0%

4.3 Dev Set (qrels.dev.tsv)

Table 4: Performance on Dev Set (1,073 binary judgments)

System	MAP	MRR@10	Recall@100
BM25	0.1478	0.1507	0.3965
HNSW Dense	0.0108	0.0108	0.9321
Hybrid (top 1000)	0.3137	0.3176	0.4195
Improvement	+112.2%	+110.8%	+5.8%

4.4 Key Observations

1. **Hybrid Dominates Precision Metrics:** The hybrid system achieves 82-90% MRR@10 across all test sets, far exceeding both BM25 (56-59%) and HNSW (27-29%). This represents a 47-112% improvement over BM25 baseline.
2. **NDCG Performance Validates Ranking Quality:** Hybrid achieves 48-59% NDCG@10 and 43-48% NDCG@100, demonstrating superior ranking of graded relevance. The

73-81% improvement in NDCG@10 over BM25 indicates the hybrid system excels at placing highly relevant documents at top positions.

3. **HNSW Excels at Recall:** Pure dense retrieval achieves 93% recall on the dev set and 55-61% on TREC DL sets, finding nearly all relevant documents. However, this comes at the cost of poor precision (1-2% MRR@10).
4. **Precision-Recall Trade-off is Clear:** HNSW achieves 69-93% recall but only 1-29% MRR@10, while hybrid balances with 40-42% recall and 82-90% MRR@10. BM25 falls in the middle with 34-40% recall and 15-59% MRR@10.
5. **Hybrid Improves BM25 Recall:** By using top-1000 BM25 candidates instead of top-100, hybrid recall improved from 39.45% to 41.95% on the dev set, a 6.3% gain while maintaining superior precision.
6. **Dev Set Shows Largest Gains:** MAP more than doubled (14.78% \rightarrow 31.37%), and MRR@10 improved by 111%, suggesting hybrid is especially effective for binary relevance tasks where lexical matching is less reliable.
7. **Dense Retrieval Fails at Top-10 Precision:** Despite having the best NDCG@100 scores (36-42%), HNSW’s NDCG@10 is only 9-15%, indicating poor ranking at top positions despite good overall relevance coverage.

5 Analysis and Discussion

5.1 Trade-offs in Retrieval Quality

5.1.1 Precision vs. Recall

Our results reveal a fundamental precision-recall trade-off across the three systems:

- **BM25:** Achieves moderate precision (15-59% MRR@10) and recall (34-40%), providing a balanced baseline with minimal computational cost
- **HNSW:** Exceptional recall (55-93%) but poor precision (1-29% MRR@10), useful for high-recall applications but not production search
- **Hybrid:** Best precision (32-90% MRR@10) with reasonable recall (40-42%), optimal for user-facing search applications

Example Scenario:

- Query: "What causes high blood pressure?"
- BM25 finds: Documents with exact terms "high blood pressure"
- HNSW finds: Documents about "hypertension factors" (semantic match)
- Hybrid: Retrieves BM25 candidates, re-ranks to promote semantic matches

5.1.2 When Each System Excels

Table 5: System strengths by query type

Query Type	Best System	Reason
Named entities	BM25	Exact matching crucial
Conceptual	HNSW/Hybrid	Semantic understanding needed
Multi-term precise	BM25	Boolean AND logic effective
Paraphrased	HNSW/Hybrid	Vocabulary mismatch common
Rare terms	BM25	IDF weighting powerful
Common concepts	Hybrid	Disambiguation via embeddings

5.2 Efficiency Analysis

5.2.1 Time Complexity

Table 6: Computational complexity comparison

Operation	BM25	HNSW	Hybrid
Index construction	$O(N \log N)$	$O(N \log N)$	-
Query time (single)	$O(q \cdot k)$	$O(\log N)$	$O(q \cdot k + K \log K)$
Space	$O(V + N)$	$O(N \cdot d)$	$O(V + N + K \cdot d)$

Where:

- N = number of documents (1M)
- V = vocabulary size (445K terms)
- d = embedding dimension (384)
- $|q|$ = query length
- k = average postings list length
- K = candidate set size (1000)

5.2.2 Memory Usage

Table 7: Memory footprint comparison

Component	Size	System
BM25 lexicon + metadata	195 MB	BM25
BM25 inverted index (disk)	98 MB	BM25
Passage embeddings	771 MB	HNSW, Hybrid
Query embeddings	156 MB	HNSW, Hybrid
Total BM25	195 MB	-
Total HNSW	927 MB	-
Total Hybrid	927 MB + BM25 disk	-

Key Insight: BM25 uses $4.7\times$ less memory than dense retrieval, making it more scalable for resource-constrained environments.

5.2.3 Query Latency

Table 8: Average query latency

System	Single Query	Batch (202K queries)
BM25	27-58 ms	~45 seconds
HNSW	12 ms (estimated)	~2.5 minutes
Hybrid (per query)	~600 ms	~2 minutes

Analysis:

- BM25 is fastest for single queries
- HNSW benefits from batch processing and SIMD optimization
- Hybrid adds re-ranking overhead but provides superior quality

6 Conclusion

6.1 Summary of Findings

This project successfully implemented and compared three information retrieval paradigms on the MS-MARCO dataset. Our key findings are:

1. **Hybrid Re-ranking Achieves Best Overall Performance:** With MRR@10 of 82-90%, NDCG@10 of 48-59%, and NDCG@100 of 43-48%, the hybrid system significantly outperforms both BM25 and pure dense retrieval across all precision metrics.

2. **BM25 Remains Highly Competitive:** Despite being a 40-year-old algorithm, BM25 achieves 56-59% MRR@10 on TREC DL benchmarks with minimal computational cost (195MB RAM, 27-58ms latency), making it an excellent baseline.
3. **Dense Retrieval Solves the Vocabulary Mismatch Problem:** HNSW achieves 55-93% recall, successfully finding relevant documents even when query and document terms differ. However, its 1-29% MRR@10 makes it unsuitable for standalone use in production.
4. **Two-Stage Retrieval is Practical and Effective:** Using BM25 for candidate generation (fast, cheap) and dense vectors for re-ranking (precise) combines the strengths of both approaches while maintaining computational efficiency.
5. **Top-K Selection Matters:** Increasing from top-100 to top-1000 BM25 candidates improved hybrid recall from 39.45% to 41.95% (+6.3%), demonstrating that candidate pool size directly impacts final performance.
6. **NDCG Metrics Reveal Ranking Quality:** The 73-81% improvement in NDCG@10 (hybrid vs. BM25) demonstrates that semantic re-ranking excels at placing highly relevant documents at top positions, crucial for user satisfaction.

6.2 Summary

Table 9: Decision matrix for system selection

Scenario	Recommended System	
Production search	web	Hybrid (best user experience)
Low-latency applications	BM25 (27ms vs 600ms)	
Resource-constrained	BM25 (195MB vs 927MB)	
Multi-lingual search	HNSW or Hybrid (with multilingual embeddings)	
Legal/medical search	BM25 (explainability crucial)	
Exploratory research	HNSW (high recall important)	
E-discovery systems	HNSW (must find all relevant documents)	
Real-time applications	BM25 (minimal latency critical)	

6.3 Final Remarks

This project demonstrates that information retrieval remains an active area of research where classical algorithms (BM25) and modern deep learning (dense retrieval) each have distinct advantages. The future of search lies not in replacing one approach with another, but in intelligent combination of complementary techniques. Our hybrid system achieving 82-90% MRR@10 and 48-59% NDCG@10 validates this hybrid philosophy and provides a strong foundation for building production-grade search systems.

The choice of retrieval system ultimately depends on the specific application requirements, balancing retrieval quality, computational cost, memory constraints, and latency requirements. By understanding these trade-offs, practitioners can make informed decisions when designing information retrieval systems.